

code

September 12, 2021

```
[21]: import pandas as pd
      from matplotlib import pyplot as plt
      import numpy as np
```

```
[22]: df = pd.read_csv("./FIFA-21 Complete.csv")
      df.head()
```

```
[22]:
```

	player_id	name	nationality	position	overall	age	hits	\
0	158023	Lionel Messi	Argentina	ST CF RW	94	33	299	
1	20801	Cristiano Ronaldo	Portugal	ST LW	93	35	276	
2	190871	Neymar Jr	Brazil	CAM LW	92	28	186	
3	203376	Virgil van Dijk	Netherlands	CB	91	29	127	
4	200389	Jan Oblak	Slovenia	GK	91	27	47	

	potential	team
0	94	FC Barcelona
1	93	Juventus
2	92	Paris Saint-Germain
3	92	Liverpool
4	93	Atlético Madrid

```
[23]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17981 entries, 0 to 17980
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   player_id       17981 non-null  int64
1   name            17981 non-null  object
2   nationality      17981 non-null  object
3   position        17981 non-null  object
4   overall         17981 non-null  int64
5   age             17981 non-null  int64
6   hits            17981 non-null  int64
7   potential       17981 non-null  int64
8   team            17981 non-null  object
```

```
dtypes: int64(5), object(4)
memory usage: 1.2+ MB
```

```
[24]: X = df[["overall", "age", "hits", "potential"]]
```

```
[25]: from sklearn.preprocessing import MinMaxScaler
X_sca = MinMaxScaler()
X = X_sca.fit_transform(X)
```

```
[26]: dfq = pd.DataFrame(data=X, columns=["overall", "age", "hits", "potential"])
```

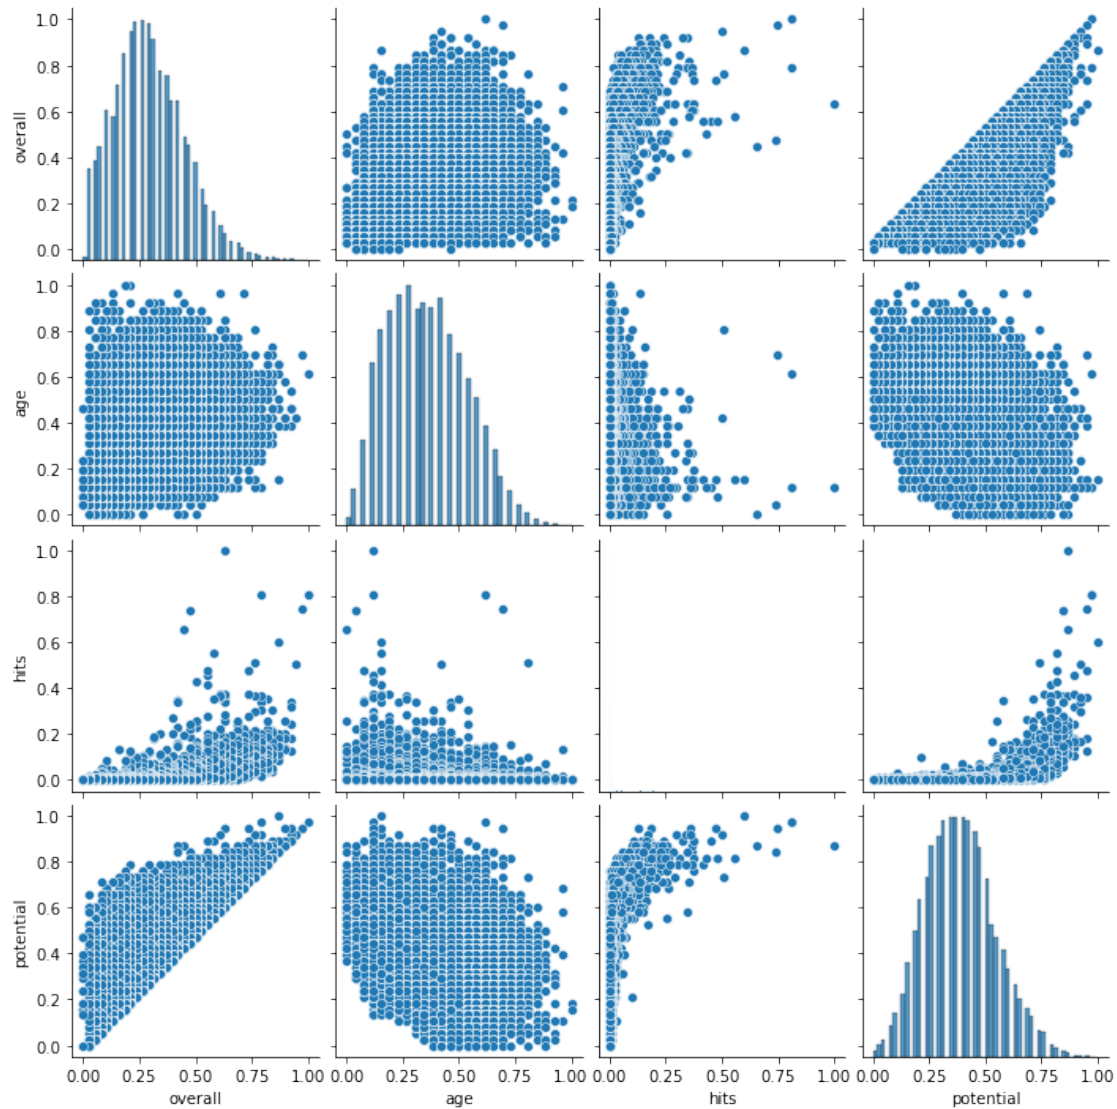
```
[27]: dfq.describe()
```

```
[27]:
```

	overall	age	hits	potential
count	17981.000000	17981.000000	17981.000000	17981.000000
mean	0.296693	0.358132	0.007249	0.387844
std	0.155905	0.175234	0.029235	0.156894
min	0.000000	0.000000	0.000000	0.000000
25%	0.184211	0.230769	0.000000	0.263158
50%	0.289474	0.346154	0.000000	0.368421
75%	0.394737	0.500000	0.005391	0.500000
max	1.000000	1.000000	1.000000	1.000000

```
[28]: import seaborn as sns
sns.pairplot(dfq)
```

```
[28]: <seaborn.axisgrid.PairGrid at 0x7fdd880bb850>
```



```
[29]: X = np.array(dfq["hits"])
      y = np.array(dfq["potential"])
      X = X.reshape(-1, 1)
      y = y.reshape(-1, 1)
      print(X.shape, y.shape)
```

```
(17981, 1) (17981, 1)
```

```
[30]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(
          X, y, test_size=0.20, random_state=0)
```

```
[31]: from sklearn.linear_model import LinearRegression
reg = LinearRegression()
print(X_train.shape, y_train.shape)
reg.fit(X_train, y_train)
```

(14384, 1) (14384, 1)


```
[31]: LinearRegression()
```

```
[32]: y_pred = reg.predict(X_test)
```

```
[33]: reg.score(X_test, y_test)
```

```
[33]: 0.16111153706859516
```

```
[34]: plt.scatter(X, y)
plt.plot(X_train, reg.predict(X_train), color="red")
plt.show()
```



```
[35]: from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
kf = KFold(n_splits=5)
model = LinearRegression()
scores = cross_val_score(model, X_train, y_train, scoring='r2', cv=kf)

print("Avg accuracy: {}".format(scores.mean()))
```

Avg accuracy: 0.18585985402502103