

Assignment 3

Drug Clasification

Importing Libraries

```
In [5]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.graph_objs as go
import plotly.express as px
import warnings
warnings.simplefilter("ignore")
```

Loading up the data

```
In [6]: data = pd.read_csv("drug200.csv")
data.head()
```

```
Out[6]:
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	DrugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	DrugY

- **Age:** Age of the patient
- **Sex:** Gender of the patients
- **BP:** Blood Pressure of the patient
- **Cholesterol:** Cholesterol of the patient

- **Na_to_K**: Sodium to Potassium ratio in patient's blood
- **Drug**: Drug type give to patients

```
In [7]: # Looking for missing values in the dataset
data.isna().sum()
```

```
Out[7]: Age          0
Sex          0
BP           0
Cholesterol  0
Na_to_K      0
Drug         0
dtype: int64
```

```
In [8]: data.shape
```

```
Out[8]: (200, 6)
```

```
In [9]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age             200 non-null   int64
1   Sex             200 non-null   object
2   BP              200 non-null   object
3   Cholesterol      200 non-null   object
4   Na_to_K         200 non-null   float64
5   Drug            200 non-null   object
dtypes: float64(1), int64(1), object(4)
memory usage: 9.5+ KB
```


```
In [10]: data.dtypes
```

```
Out[10]: Age          int64
Sex          object
BP           object
Cholesterol  object
Na_to_K      float64
Drug         object
dtype: object
```

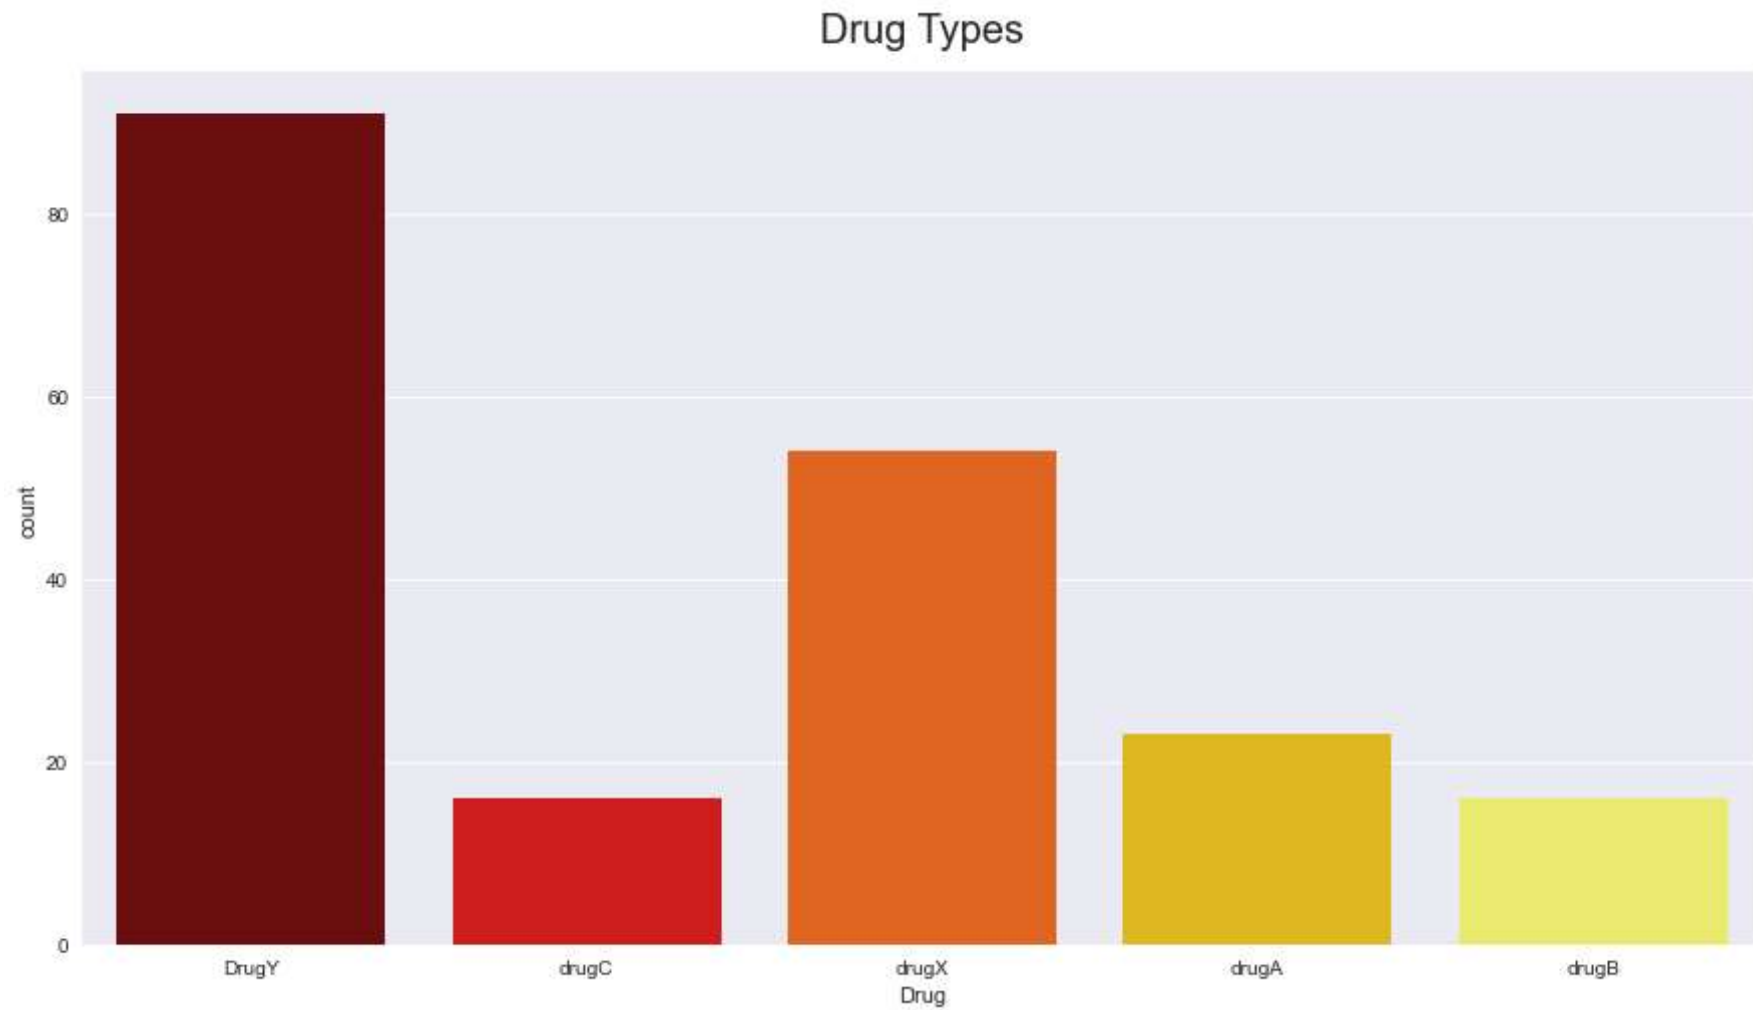
```
In [11]: data['Drug'].value_counts()
```

```
Out[11]: DrugY    91  
         drugX    54  
         drugA    23  
         drugB    16  
         drugC    16  
         Name: Drug, dtype: int64
```

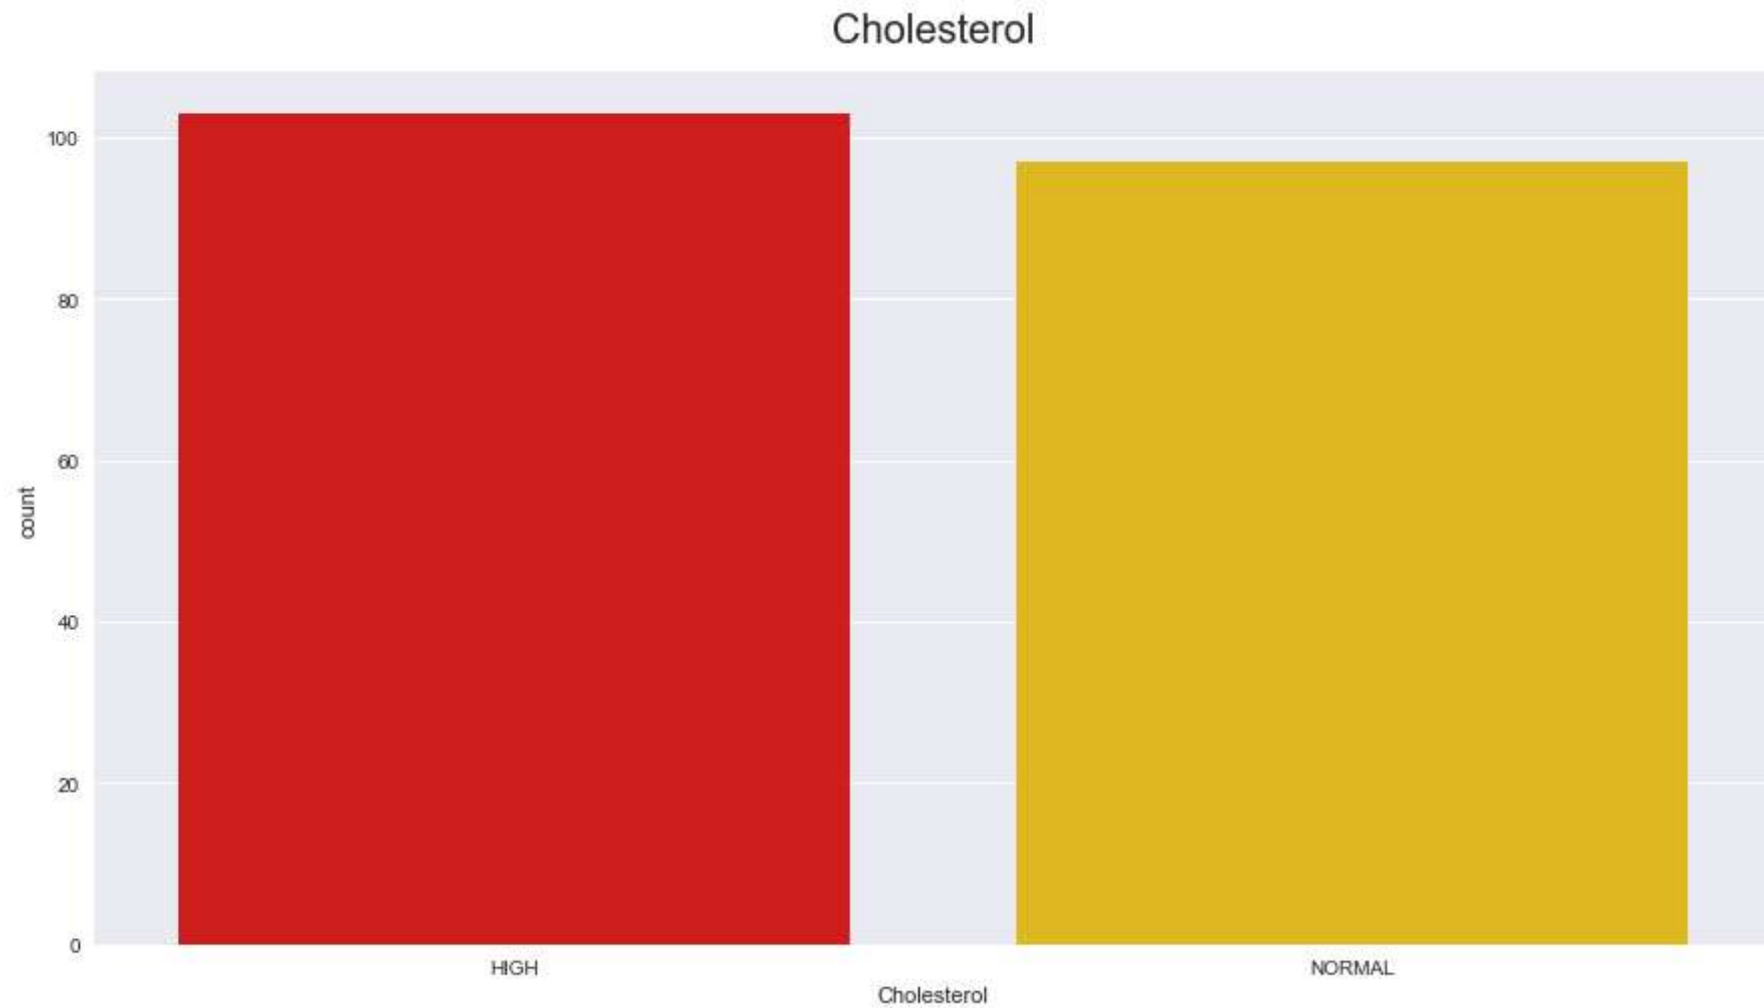
```
In [12]: plt.style.use("seaborn")  
         plt.figure(figsize=(15,8))  
         plt.title("Genders", fontsize=20, y=1.02)  
         sns.countplot(x = data.Sex, palette="hot")  
         plt.show()
```



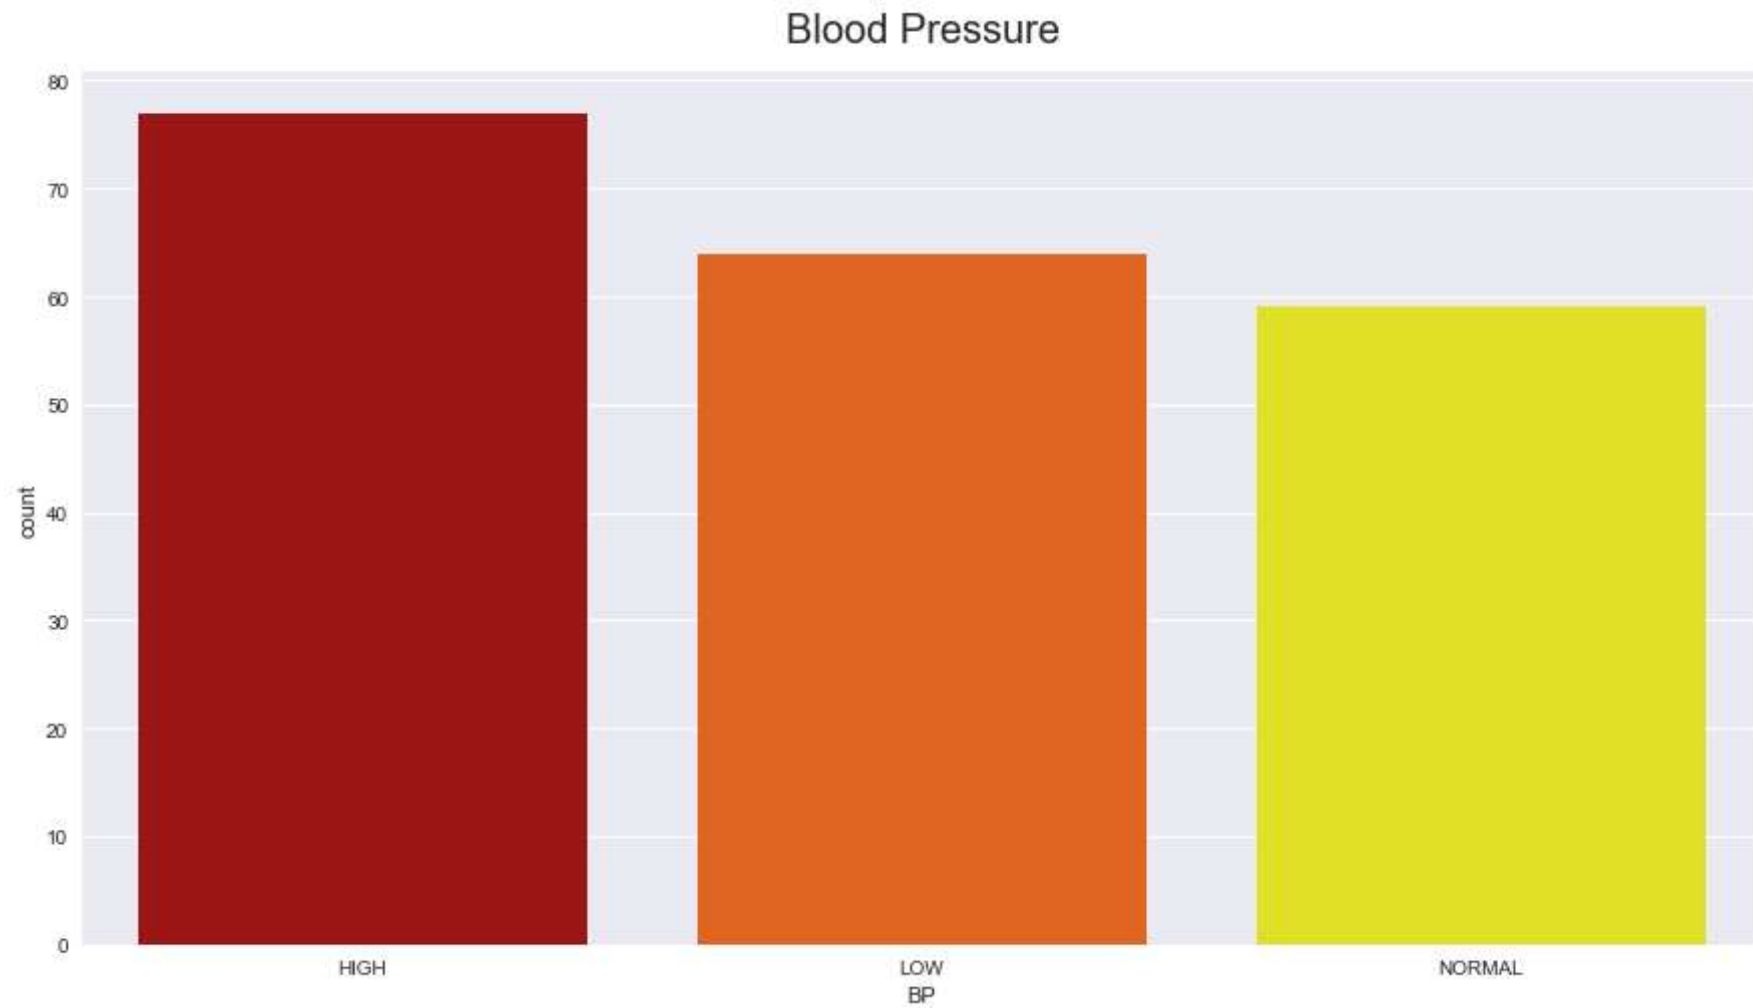
```
In [13]: plt.figure(figsize=(15,8))  
plt.title("Drug Types", fontsize=20, y=1.02)  
sns.countplot(x = data.Drug, palette="hot")  
plt.show()
```




```
In [14]: plt.figure(figsize=(15,8))  
plt.title("Cholesterol", fontsize=20, y=1.02)  
sns.countplot(x = data.Cholesterol, palette="hot")  
plt.show()
```



```
In [15]: plt.figure(figsize=(15,8))  
plt.title("Blood Pressure", fontsize=20, y=1.02)  
sns.countplot(x = data.BP, palette="hot")  
plt.show()
```



```
In [16]: plt.style.use("seaborn")
fig, ax = plt.subplots(figsize=(15,8))
sns.histplot(data["Age"], kde=True, bins=25, color="firebrick")
plt.title("Age of the patients", fontsize=20, y=1.02)
ax.set_xlabel("Age", fontsize=15);
```



```
In [17]: data.head(1)
```


```
Out[17]:
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	DrugY

```
In [18]: plt.style.use("seaborn")
fig, ax = plt.subplots(figsize=(15,8))
sns.histplot(data["Na_to_K"], color="red", kde=True, bins=25)
```





```
plt.title("Sodium to Potassium ratio in patient's blood", fontsize=20, y=1.02)
ax.set_xlabel("Na_to_K", fontsize=15);
```



```
In [19]: plt.style.use("seaborn")
fig, ax = plt.subplots(1,2, figsize=(20,8))

sns.barplot(x = "Sex", y = "Count", hue = "Cholesterol", data = data.groupby(["Sex", "Cholesterol"]).size().reset_index(name = "Count"))
ax[0].set_xlabel("Sex", fontsize=14);



sns.barplot(x = "Sex", y = "Count", hue = "BP", data = data.groupby(["Sex", "BP"]).size().reset_index(name = "Count"), palette="hsv")
ax[1].set_xlabel("Sex", fontsize=14);
```



```
In [20]: plt.style.use("seaborn")
fig, ax = plt.subplots(1,2, figsize=(20,8))

sns.barplot(x = "BP", y = "Count", hue = "Cholesterol", data = data.groupby(["BP", "Cholesterol"]).size().reset_index(name = "Count"))
ax[0].set_xlabel("BP", fontsize=14);



sns.barplot(x = "Drug", y = "Count", hue = "Cholesterol", data = data.groupby(["Drug", "Cholesterol"]).size().reset_index(name = "Count"))
ax[1].set_xlabel("Drug", fontsize=14);
```



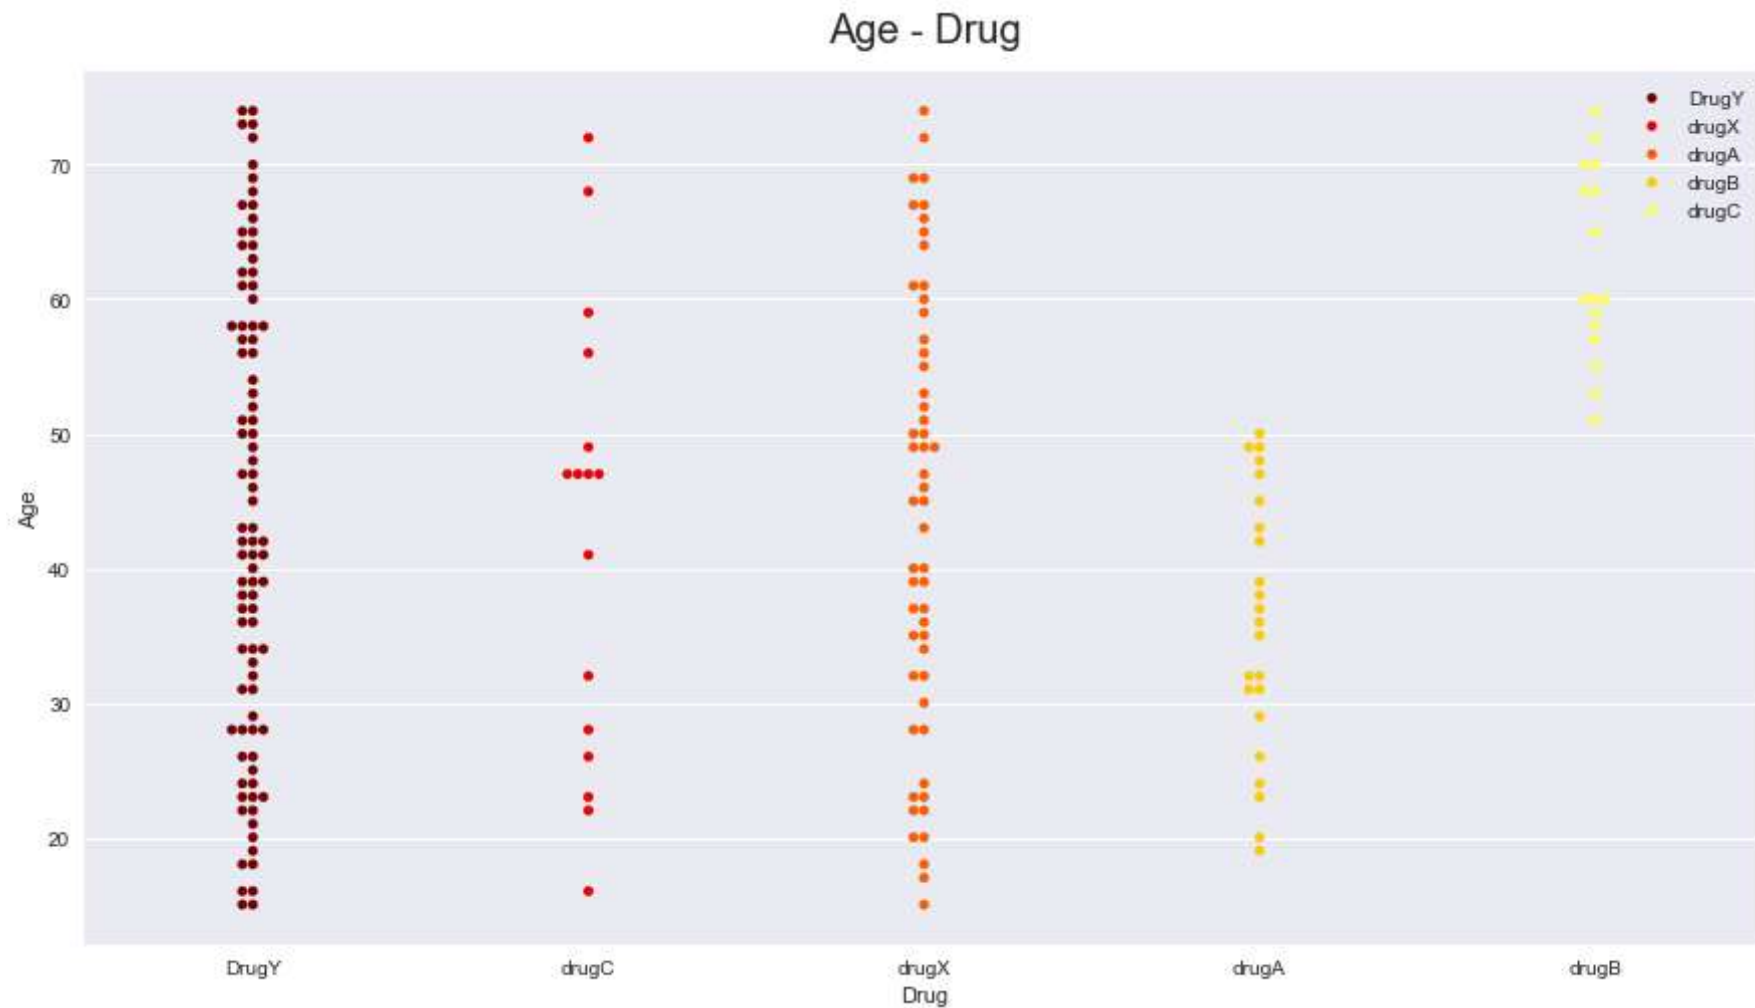
```
In [21]: plt.style.use("seaborn")
fig, ax = plt.subplots(1,2, figsize=(20,8))

sns.barplot(x = "Sex", y = "Count", hue = "Drug", data = data.groupby(["Sex", "Drug"]).size().reset_index(name = "Count"), palette=
ax[0].set_xlabel("Sex",fontsize=14);

sns.barplot(x = "Drug", y = "Count", hue = "BP", data = data.groupby(["Drug", "BP"]).size().reset_index(name = "Count"), palette="
ax[1].set_xlabel("Drug",fontsize=14);
```



```
In [22]: plt.figure(figsize = (15,8))
sns.swarmplot(x = "Drug", y = "Age", data = data, palette="hot")
plt.legend(data.Drug.value_counts().index)
plt.title("Age - Drug", fontsize=20, y=1.02)
plt.show()
```



```
In [23]: data.dtypes
```

```
Out[23]: Age          int64  
Sex          object  
BP           object  
Cholesterol  object  
Na_to_K      float64  
Drug         object  
dtype: object
```

```
In [24]: # Converting the non-numeric values into numeric values
```

```
data['Sex'] = data['Sex'].map({'M': 1, 'F': 2})
data['BP'] = data['BP'].map({'HIGH': 1, "NORMAL" : 2, "LOW" : 3})
data['Cholesterol'] = data['Cholesterol'].map({'HIGH': 1, "NORMAL" : 2})
data["Drug"] = data["Drug"].map({"DrugY":1, "drugC":2, "drugX":3, "drugA":4, "drugB":5})
```

In [25]: `data.head()`

Out[25]:

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	2	1	1	25.355	1
1	47	1	3	1	13.093	2
2	47	1	3	1	10.114	2
3	28	2	2	1	7.798	3
4	61	2	3	1	18.043	1

Splitting the data into training and test datasets

Here, we are trying to predict the Drug type that is to be prescribed to the patient using the given data. Hence, the "Drug Type" will be the y label and rest of the data will be the X or the input data.

In [26]:

```
# X data
X = data.drop("Drug", axis=1)
X.head()
```

Out[26]:

	Age	Sex	BP	Cholesterol	Na_to_K
0	23	2	1	1	25.355
1	47	1	3	1	13.093
2	47	1	3	1	10.114
3	28	2	2	1	7.798
4	61	2	3	1	18.043

In [27]:

```
# y data
y = data["Drug"]
```

```
y.head()
```

```
Out[27]: 0    1  
1    2  
2    2  
3    3  
4    1  
Name: Drug, dtype: int64
```

```
In [28]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [29]: len(X_train), len(X_test)
```

```
Out[29]: (160, 40)
```

```
In [30]: # Scaling the data  
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

Logistic Regression

```
In [31]: from sklearn.linear_model import LogisticRegression  
lr = LogisticRegression()  
lr.fit(X_train, y_train)
```


```
Out[31]: LogisticRegression()
```

```
In [32]: LogisticRegressionScore = lr.score(X_test, y_test)  
print("Accuracy obtained by Logistic Regression model:", LogisticRegressionScore*100)
```

```
Accuracy obtained by Logistic Regression model: 92.5
```

```
In [49]: # Having a Look at the confusion matrix for Logistic Regression  
  
from sklearn.metrics import confusion_matrix, classification_report  
  
y_pred_lr = lr.predict(X_test)  
cf_matrix = confusion_matrix(y_test, y_pred_lr)
```

```
sns.heatmap(cf_matrix, annot=True, cmap="vlag_r")
plt.title("Confusion Matrix for Logistic Regression", fontsize=14, fontname="Helvetica", y=1.03);
```



In [34]: *# Having a Look at the classification report of Logistic Regression*

```
from sklearn import metrics
print(metrics.classification_report(y_test, y_pred_lr))
```

	precision	recall	f1-score	support
1	1.00	0.93	0.97	15
2	1.00	0.60	0.75	5
3	0.85	1.00	0.92	11
4	0.86	1.00	0.92	6
5	1.00	1.00	1.00	3
accuracy			0.93	40
macro avg	0.94	0.91	0.91	40
weighted avg	0.94	0.93	0.92	40

Random Forest Classifier

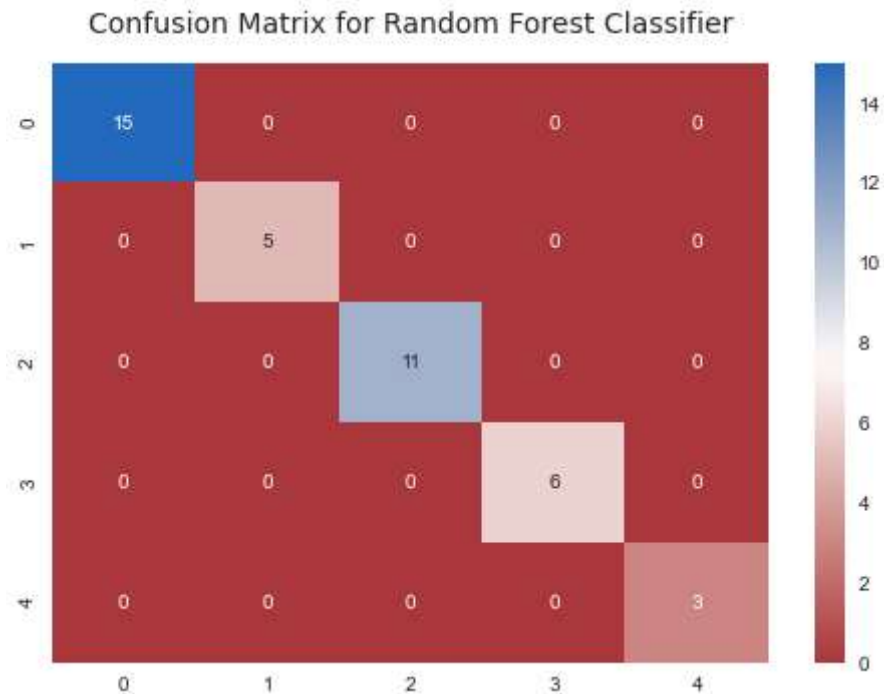
```
In [35]: from sklearn.ensemble import RandomForestClassifier  
rfc = RandomForestClassifier()  
rfc.fit(X_train, y_train)
```

Out[35]: RandomForestClassifier()

```
In [36]: RandomForestClassifierScore = rfc.score(X_test,y_test)  
print("Accacy obtained by Random Forest Classifier :", RandomForestClassifierScore*100)
```

Accacy obtained by Random Forest Classifier : 100.0

```
In [37]: # Confusion Matrix of Random Forest Classifier  
  
y_pred_rfc = rfc.predict(X_test)  
cf_matrix = confusion_matrix(y_test, y_pred_rfc)  
sns.heatmap(cf_matrix, annot=True, cmap="vlag_r")  
plt.title("Confusion Matrix for Random Forest Classifier", fontsize=14, fontname="Helvetica", y=1.03);
```



```
In [38]: print(metrics.classification_report(y_test, y_pred_rfc))
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	15
2	1.00	1.00	1.00	5
3	1.00	1.00	1.00	11
4	1.00	1.00	1.00	6
5	1.00	1.00	1.00	3
accuracy			1.00	40
macro avg	1.00	1.00	1.00	40
weighted avg	1.00	1.00	1.00	40

K Neighbors Classifier

```
In [39]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
```


```
Out[39]: KNeighborsClassifier()
```

```
In [40]: KNeighborsClassifierScore = knn.score(X_test, y_test)
print("Accuracy obtained by K Neighbors Classifier :", KNeighborsClassifierScore*100)
```

Accuracy obtained by K Neighbors Classifier : 85.0

```
In [41]: # Confusion Matrix

y_pred_knn = knn.predict(X_test)
cf_matrix = confusion_matrix(y_test, y_pred_knn)
sns.heatmap(cf_matrix, annot=True, cmap="vlag_r")
plt.title("Confusion Matrix for K Neighbors Classifier", fontsize=14, fontname="Helvetica", y=1.03);
```



```
In [42]: print(metrics.classification_report(y_test,y_pred_knn))
```

	precision	recall	f1-score	support
1	0.92	0.80	0.86	15
2	1.00	0.80	0.89	5
3	0.77	0.91	0.83	11
4	0.71	0.83	0.77	6
5	1.00	1.00	1.00	3
accuracy			0.85	40
macro avg	0.88	0.87	0.87	40
weighted avg	0.86	0.85	0.85	40

Decision Tree Classifier


```
In [43]: from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier()
tree.fit(X_train, y_train)
```

Out[43]: DecisionTreeClassifier()

```
In [44]: DecisionTreeClassifierScore = tree.score(X_test,y_test)
print("Accuracy obtained by Decision Tree Classifier :", DecisionTreeClassifierScore*100)
```

Accuracy obtained by Decision Tree Classifier : 100.0

```
In [45]: y_pred_tree = tree.predict(X_test)
cf_matrix = confusion_matrix(y_test, y_pred_tree)
sns.heatmap(cf_matrix, annot=True, cmap="vlag_r")
plt.title("Confusion Metrix for Decision Tree Classifier", fontsize=14, fontname="Helvetica", y=1.03);
```



```
In [46]: print(metrics.classification_report(y_test, y_pred_tree));
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	15
2	1.00	1.00	1.00	5
3	1.00	1.00	1.00	11
4	1.00	1.00	1.00	6
5	1.00	1.00	1.00	3

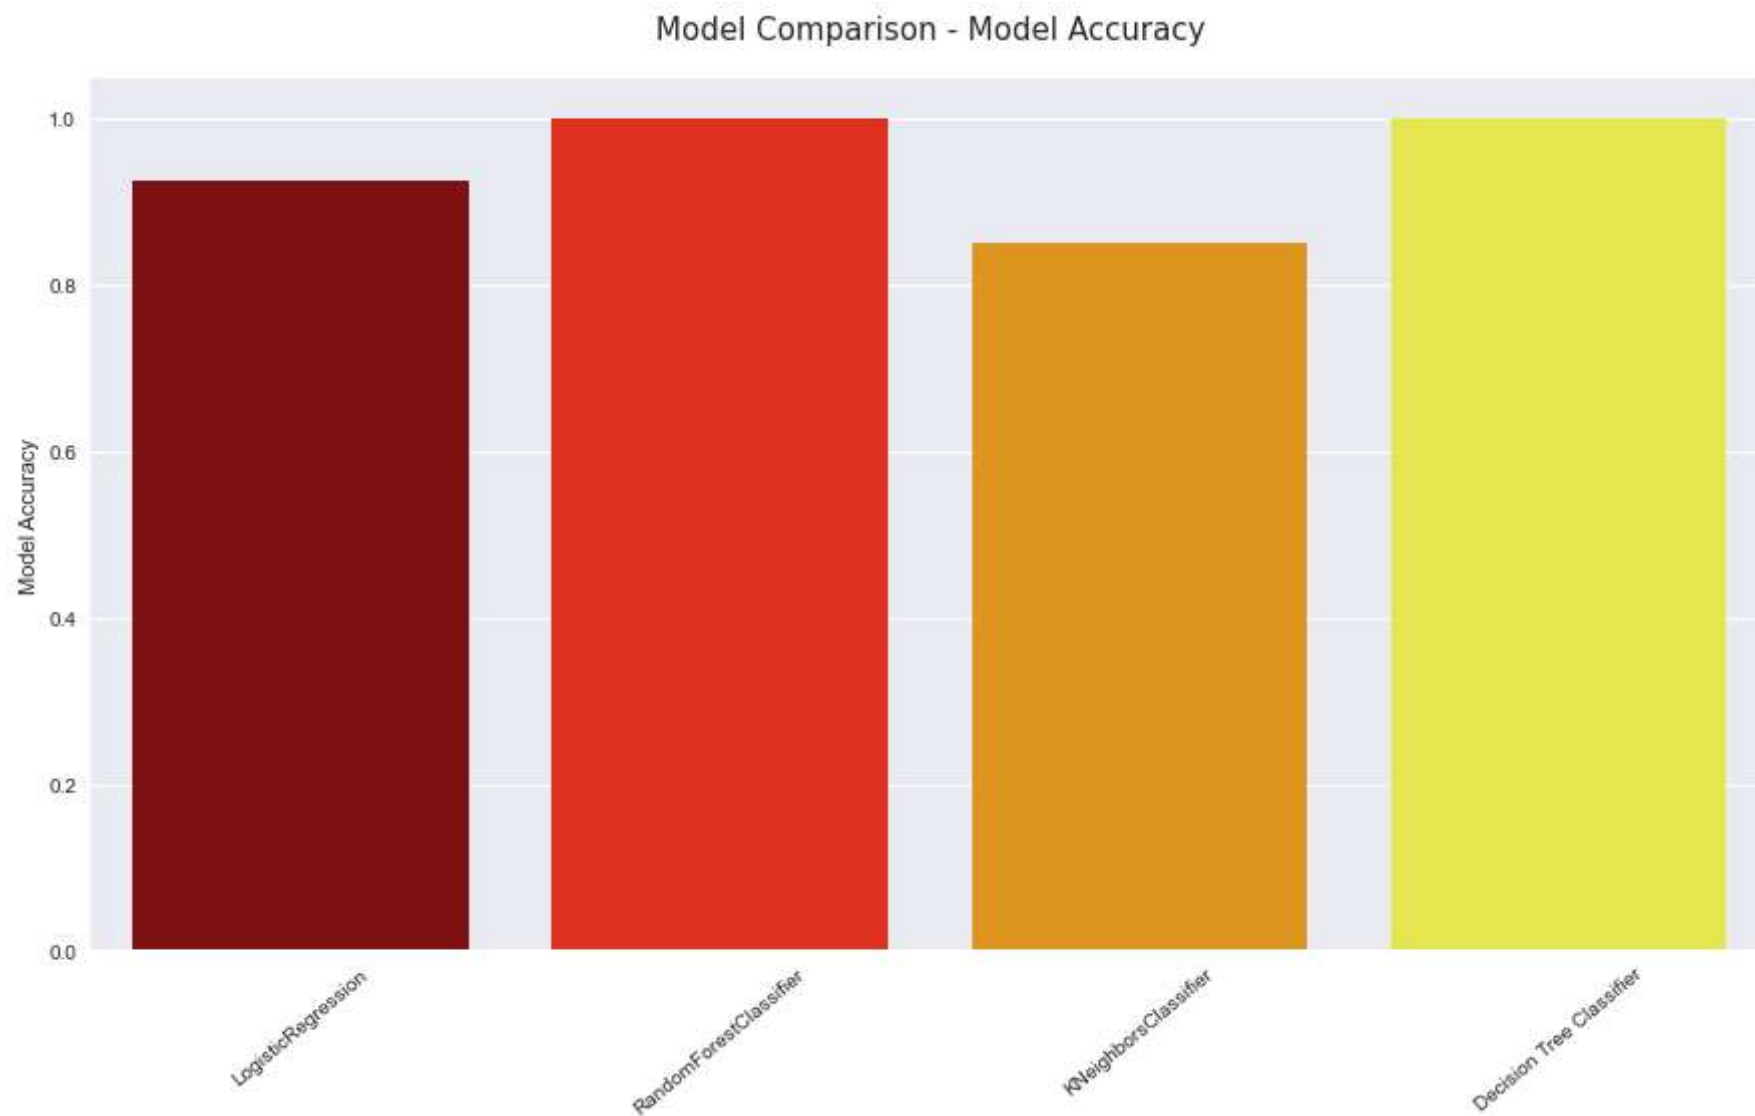
accuracy			1.00	40
macro avg	1.00	1.00	1.00	40
weighted avg	1.00	1.00	1.00	40

```
In [48]: plt.style.use("seaborn")

x = ["LogisticRegression",
     "RandomForestClassifier",
     "KNeighborsClassifier",
     "Decision Tree Classifier"]

y = [LogisticRegressionScore,
     RandomForestClassifierScore,
     KNeighborsClassifierScore,
     DecisionTreeClassifierScore]

fig, ax = plt.subplots(figsize=(15,8))
sns.barplot(x=x,y=y, palette="hot");
plt.ylabel("Model Accuracy")
plt.xticks(rotation=40)
plt.title("Model Comparison - Model Accuracy", fontsize=15, fontname="Helvetica", y=1.03);
```



In []:

In []:

In []: