

```
In [256]: %matplotlib inline

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [257]: df = pd.read_csv("transistors.csv")
df.head()
```

Out[257]:

	Year	Transistors
0	1971.875000	2.308242
1	1972.307692	3.554522
2	1974.326923	6.097562
3	1979.567308	29.163776
4	1982.307692	135.772714

```
In [259]: df.dtypes
```

Out[259]: Year float64
Transistors float64
dtype: object

```
In [260]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 94 entries, 0 to 93
Data columns (total 2 columns):
Year          94 non-null float64
Transistors    94 non-null float64
dtypes: float64(2)
memory usage: 1.5 KB
```

```
In [261]: df.describe()
```

Out[261]:

	Year	Transistors
count	94.000000	9.400000e+01
mean	2002.756170	1.789723e+06
std	10.255492	5.175470e+06
min	1971.875000	2.308242e+00
25%	1996.538462	5.803409e+03
50%	2004.230769	1.139410e+05
75%	2009.699519	1.078961e+06
max	2019.800000	3.950000e+07

```
In [263]: df["Transistors"].mean()
```

Out[263]: 1789723.3063059782

```
In [264]: df.skew()
```

Out[264]: Year -0.830479
Transistors 5.293559
dtype: float64

```
In [265]: df.kurtosis()
```

Out[265]: Year 0.860388
Transistors 33.022052
dtype: float64

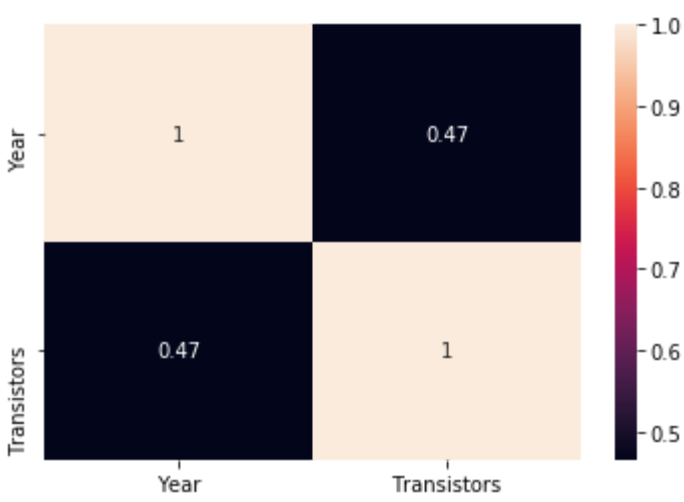
```
In [266]: df.corr()
```

Out[266]:

	Year	Transistors
Year	1.000000	0.465537
Transistors	0.465537	1.000000


```
In [267]: import seaborn as sns
sns.heatmap(df.corr(), annot=True)
```

Out[267]: <matplotlib.axes._subplots.AxesSubplot at 0x23dd711d5c0>




```
In [273]:
```

Out[273]: <matplotlib.collections.PathCollection at 0x23dd4ca85c0>



```
In [276]: df.plot.scatter("Year", "Transistors", figsize=(8,5))
```

Out[276]: <matplotlib.axes._subplots.AxesSubplot at 0x23dd74224e0>



```
In [287]: from sklearn.model_selection import train_test_split
X= df.iloc[:,0]
Y = np.log(df.iloc[:,1])
X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=1/3.0, random_state=0)
X_train = X_train.values.reshape(-1,1)
```

```
In [288]: from sklearn.linear_model import LinearRegression
```

```
regressor = LinearRegression()
regressor.fit(X_train, Y_train)
```

Out[288]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)


```
In [289]: Y_pred = regressor.predict(X_test.values.reshape(-1,1))
Y_pred
```

Out[289]: array([1.37428051, 9.80750108, 12.4551401 , 8.30390362, 7.73188284,
12.73297876, 12.74932221, 14.48695789, 9.92190524, 14.41896914,
14.31698601, 6.24462883, 12.07924073, 16.05069926, 16.83256994,
12.70029186, 6.52246749, 8.85958094, 13.89545573, 11.2947551 ,
 9.23548031, 11.08229024, 9.02301545, 5.42745629, 11.04960334,
11.73602827, 13.94984673, 14.91528705, 14.32038544, 15.65296504,
15.13964994, 3.15571664])

```
In [290]: Y_test
```

Out[290]: 2 1.807889
30 9.147379
56 12.412373
16 7.555357
13 7.447424
61 12.358406
62 13.599643
79 13.997832
33 9.983865
78 14.557448
76 14.690979
7 4.803049
51 11.629854
89 16.770421
93 17.491811
60 13.275842
8 7.096639
22 8.202959
73 13.972514
45 10.901301
26 8.175976
43 12.304439
24 9.174362
6 4.695115
42 10.550517
48 11.575887
74 14.771022
82 15.274126
77 14.648420
87 15.789592
86 15.555977
3 3.372927
Name: Transistors, dtype: float64

```
In [291]: fig = plt.figure(figsize = (16,10))
plt.scatter(X_train, Y_train)
plt.plot(np.arange(1970,2021, 1), regressor.predict(np.arange(1970,2021,1).reshape(-1,1)),color="red")
plt.title("Year vs Number of Transistors")
plt.xlabel("Year")
plt.ylabel("Number of Transistors")
plt.show()
```



```
In [292]: from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
```

```
In [294]: k=5
kf =KFold(n_splits=k, random_state=None)
result = cross_val_score(regressor, X_train, Y_train, cv =kf)
print("Avg accuracy : {}".format(result.mean()))
```

Avg accuracy : 0.9684925935287408

```
In [ ]:
```