

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
plt.style.use('fivethirtyeight')
```

```
In [4]: df = pd.read_csv('C:/Users/nisho/Documents/SEM 5/ML and core applications/winequality-red.csv')
df.head(10)
```

Out[4]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	5
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4	5
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0	7
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5	7
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5	5


## EDA

In [5]: `df.describe().T`

Out[5]:

	count	mean	std	min	25%	50%	75%	max
<b>fixed acidity</b>	1599.0	8.319637	1.741096	4.60000	7.1000	7.90000	9.200000	15.90000
<b>volatile acidity</b>	1599.0	0.527821	0.179060	0.12000	0.3900	0.52000	0.640000	1.58000
<b>citric acid</b>	1599.0	0.270976	0.194801	0.00000	0.0900	0.26000	0.420000	1.00000
<b>residual sugar</b>	1599.0	2.538806	1.409928	0.90000	1.9000	2.20000	2.600000	15.50000
<b>chlorides</b>	1599.0	0.087467	0.047065	0.01200	0.0700	0.07900	0.090000	0.61100
<b>free sulfur dioxide</b>	1599.0	15.874922	10.460157	1.00000	7.0000	14.00000	21.000000	72.00000
<b>total sulfur dioxide</b>	1599.0	46.467792	32.895324	6.00000	22.0000	38.00000	62.000000	289.00000
<b>density</b>	1599.0	0.996747	0.001887	0.99007	0.9956	0.99675	0.997835	1.00369
<b>pH</b>	1599.0	3.311113	0.154386	2.74000	3.2100	3.31000	3.400000	4.01000
<b>sulphates</b>	1599.0	0.658149	0.169507	0.33000	0.5500	0.62000	0.730000	2.00000
<b>alcohol</b>	1599.0	10.422983	1.065668	8.40000	9.5000	10.20000	11.100000	14.90000
<b>quality</b>	1599.0	5.636023	0.807569	3.00000	5.0000	6.00000	6.000000	8.00000

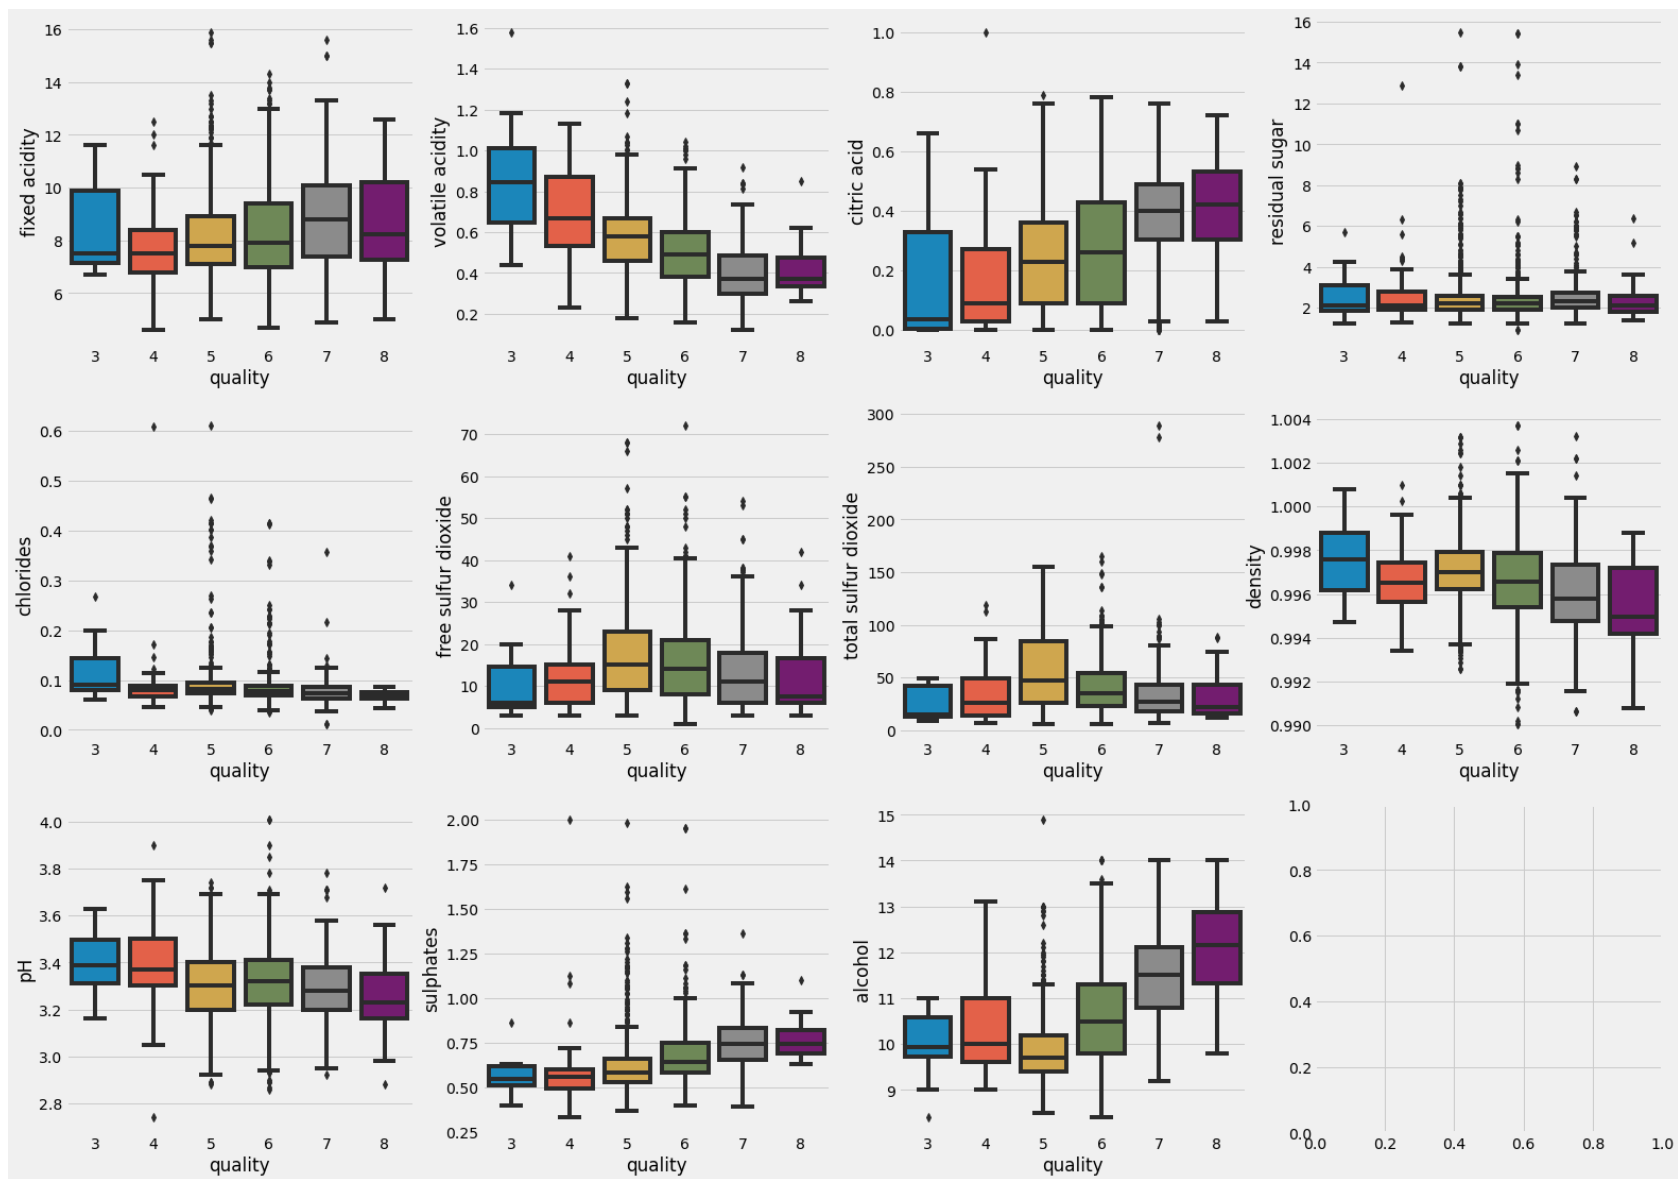
```
In [6]: sns.countplot(x='quality', data=df)  
plt.show()
```




```
In [7]: cols = list(df.columns)
fig, ax = plt.subplots(3,4, figsize=(24,18))


for i in range(11):
    j = i // 4
    k = i % 4
    sns.boxplot(y=cols[i], x = 'quality', data=df, ax = ax[j][k])

plt.show()
```



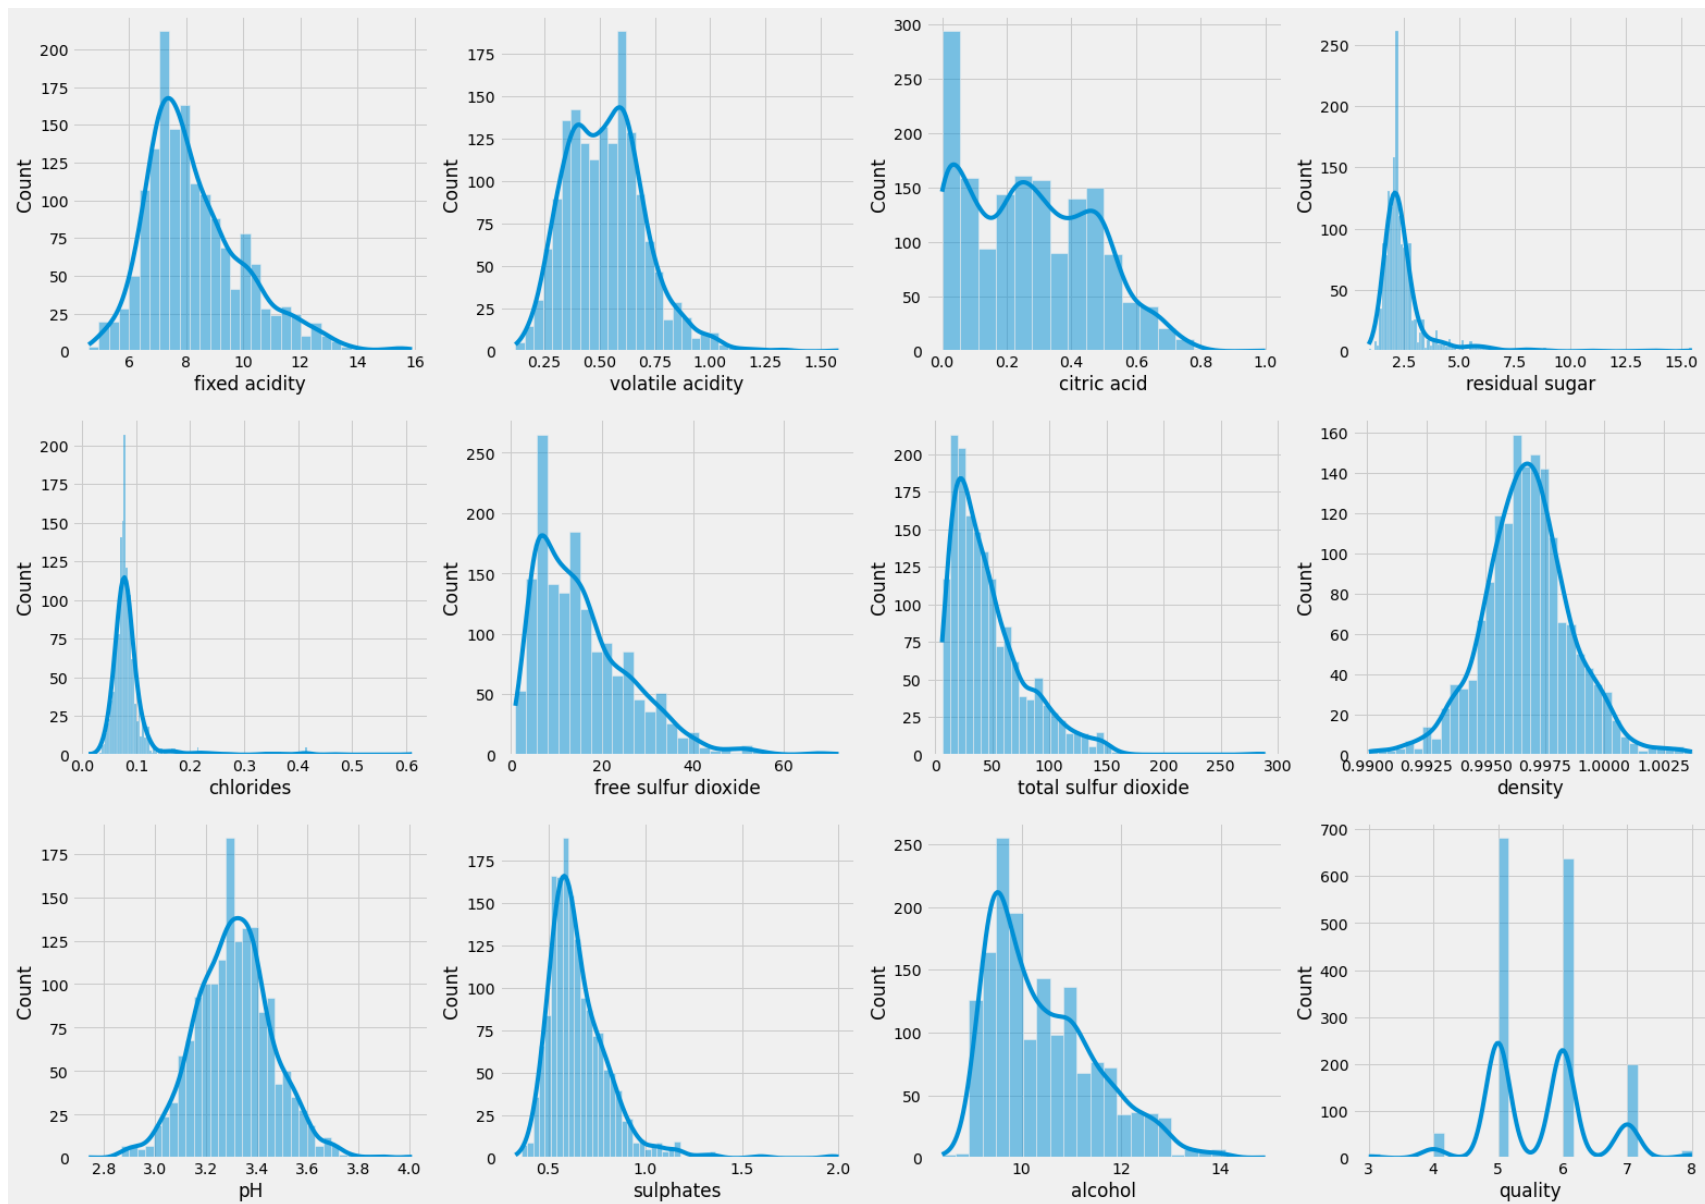
```
In [8]: corr = df.corr()
plt.figure(figsize=(12,12))
sns.heatmap(corr, cmap='YlOrRd', annot=True)
plt.show()
```





```
In [9]: fig, ax = plt.subplots(3,4, figsize=(24,18))

for i in range(12):
    j = i//4
    k = i%4
    sns.histplot(x=cols[i], data=df, ax=ax[j][k], kde=True)
```

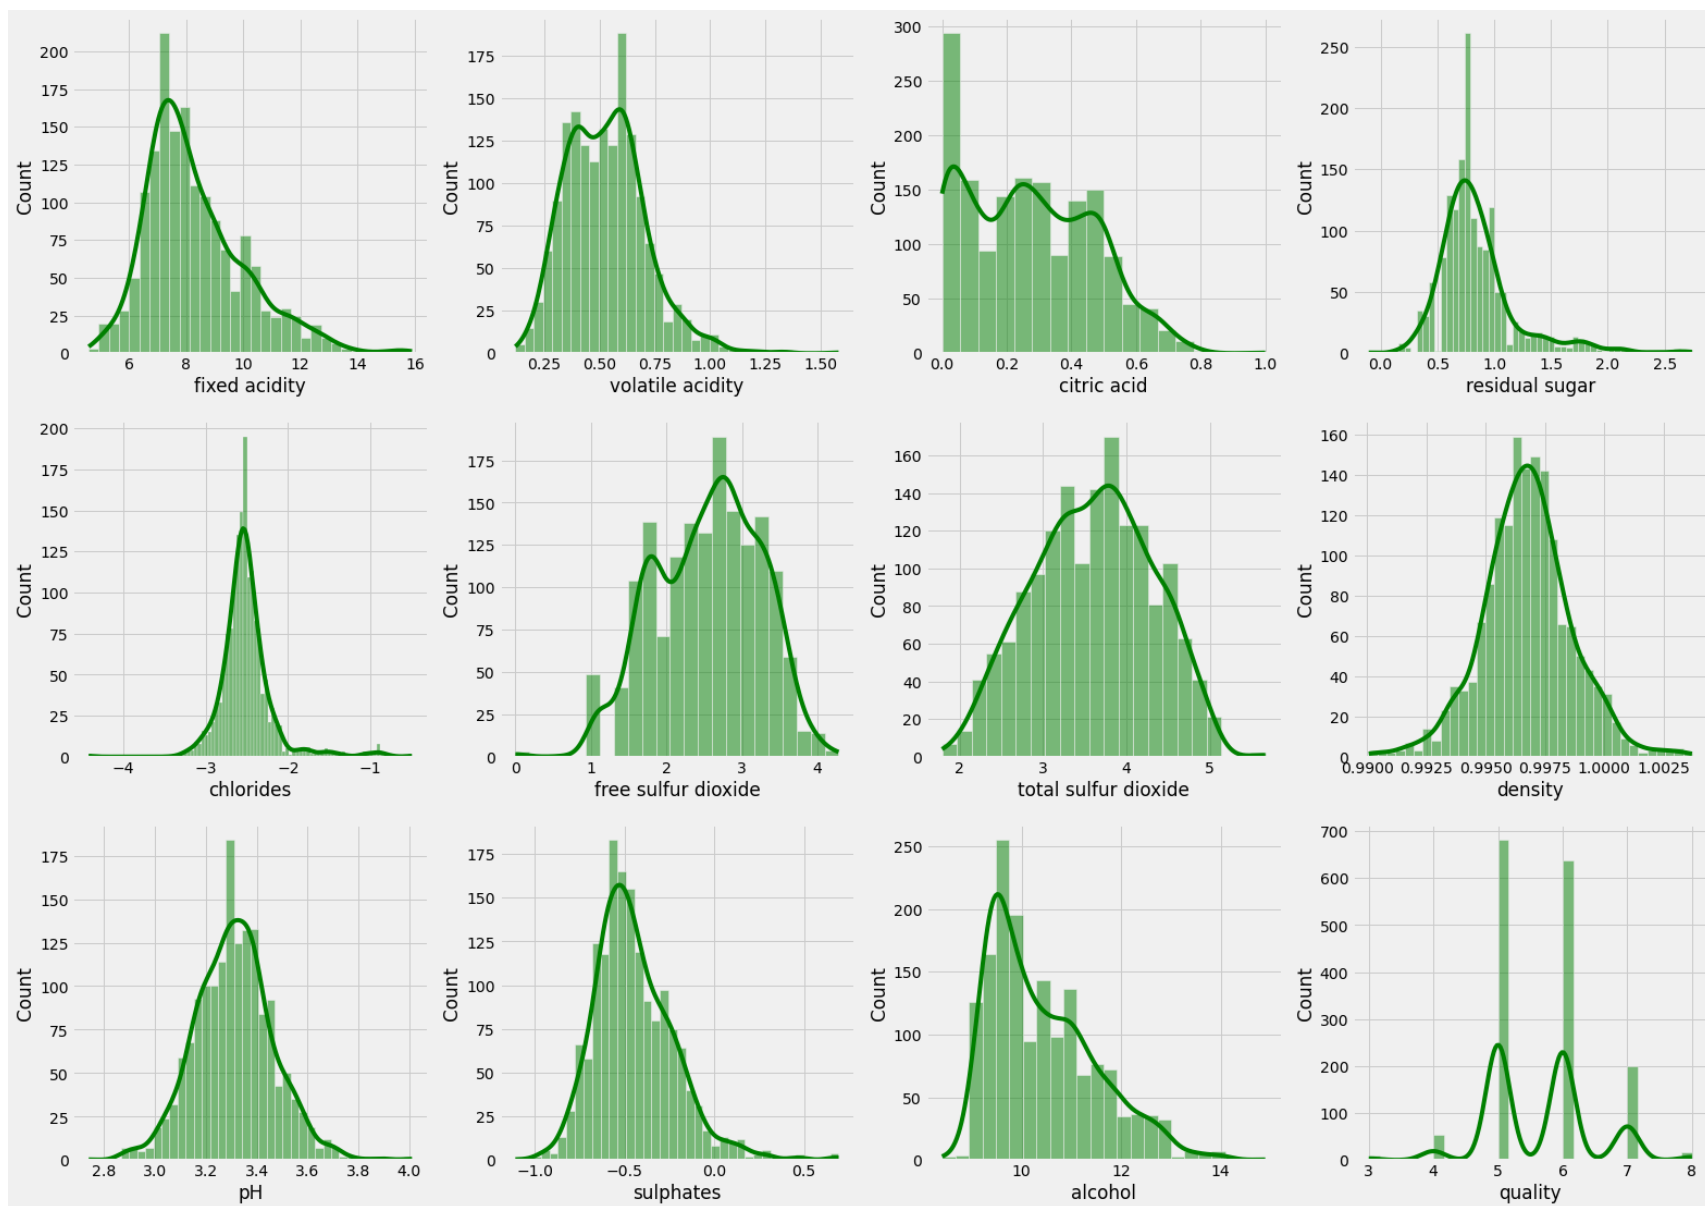




```
In [16]: skew_cols = ['residual sugar', 'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'sulphates']  
for col in skew_cols:  
    df[col] = df[col].apply(lambda x: np.log(x))
```

```
In [19]: fig, ax = plt.subplots(3,4, figsize=(24,18))

for i in range(12):
    j = i//4
    k = i%4
    sns.histplot(x=cols[i], data=df, ax=ax[j][k], kde=True, color='green')
```



```
In [20]: df_3 = df[df.quality==3]
df_4 = df[df.quality==4]
df_5 = df[df.quality==5]
df_6 = df[df.quality==6]
df_7 = df[df.quality==7]
df_8 = df[df.quality==8]
```

```
In [21]: from sklearn.utils import resample


df_3_upsampled = resample(df_3, replace=True, n_samples=400, random_state=42)
df_4_upsampled = resample(df_4, replace=True, n_samples=400, random_state=42)
df_7_upsampled = resample(df_7, replace=True, n_samples=400, random_state=42)
df_8_upsampled = resample(df_8, replace=True, n_samples=400, random_state=42)

df_5_downsampled = df_5.sample(n=400).reset_index(drop=True)
df_6_downsampled = df_6.sample(n=400).reset_index(drop=True)
```

```
In [22]: df_resampled = pd.concat([df_3_upsampled, df_4_upsampled, df_7_upsampled, df_8_upsampled,
                                   df_5_downsampled, df_6_downsampled]).reset_index(drop=True)
df_resampled.quality.value_counts().sort_index()
```


```
Out[22]: 3    400
4    400
5    400
6    400
7    400
8    400
Name: quality, dtype: int64
```


```
In [23]: sns.countplot(x='quality', data=df_resampled)  
plt.show()
```



Now we have equal sample sizes for all classes

```
In [24]: corr_2 = df_resampled.corr()
plt.figure(figsize=(12,12))
sns.heatmap(corr_2, cmap='YlOrRd', annot=True)
plt.show()
```





```
In [25]: corr_2.loc[(corr_2.quality >= 0.05) | (corr_2.quality <= -0.05), 'quality']
```

```
Out[25]: fixed acidity      0.126471
volatile acidity    -0.611091
citric acid         0.387725
chlorides          -0.365112
free sulfur dioxide  0.091006
total sulfur dioxide 0.080836
density            -0.275677
pH                 -0.271249
sulphates          0.497583
alcohol            0.580071
quality            1.000000
Name: quality, dtype: float64
```

```
In [26]: X = df_resampled.drop(['residual sugar', 'quality'], axis=1)
y = df_resampled['quality']
```

## Modelling

### decision trees

```
In [27]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier
```

```
In [28]: model = DecisionTreeClassifier(random_state = 42)
score = cross_val_score(model, X, y, cv=5)
print('Initial Score for DT classifier: ', score, '\nMean score: ', score.mean())
```

```
Initial Score for DT classifier:  [0.84791667 0.81666667 0.85208333 0.81666667 0.82083333]
Mean score:  0.8308333333333333
```

```
In [29]: from sklearn.model_selection import GridSearchCV
```

```
In [30]: model = DecisionTreeClassifier(random_state=42)
parameters = {'max_depth': [5, 10, 15, 20], 'max_features' : ['auto', 'sqrt', 'log2']}
cv = GridSearchCV(model, parameters, cv=5)
cv.fit(X, y)
```

```
Out[30]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(random_state=42),
                    param_grid={'max_depth': [5, 10, 15, 20],
                                'max_features': ['auto', 'sqrt', 'log2']})
```

```
In [31]: cv.best_score_
```

```
Out[31]: 0.8324999999999999
```

```
In [32]: cv.best_estimator_
```

```
Out[32]: DecisionTreeClassifier(max_depth=15, max_features='auto', random_state=42)
```

## Random forest

```
In [33]: from sklearn.ensemble import RandomForestClassifier
```

```
In [34]: model = RandomForestClassifier(random_state=42)

score = cross_val_score(model, X, y, cv=5, scoring='accuracy')
print('Initial score for RF: ', score, '\nMean Score: ', score.mean())
```

Initial score for RF: [0.88333333 0.8625 0.86875 0.85416667 0.84791667]  
Mean Score: 0.8633333333333333

```
In [35]: model = RandomForestClassifier(random_state=42, max_depth=15)

score = cross_val_score(model, X, y, cv=5, scoring='accuracy')
print('Improved score for RF: ', score, '\nMean Score: ', score.mean())
```

Improved score for RF: [0.89375 0.85625 0.85833333 0.86458333 0.85625 ]  
Mean Score: 0.8658333333333333

## Conclusion

Since random forest has a comparatively better mean score it outperforms decision trees